

ADVERSARIAL SEARCH


Abdus Salam Azad

Objectives

Learning how to act when the **other agents** are acting against us



Games



Games vs. Search

- Search – **no adversary**
 - Examples: path planning, scheduling activities
- Games – **adversary**
 - **Unpredictable opponent(s)**
 - Solution is strategy (strategy specifies move for every possible opponent reply).
 - **Time limits force an *approximate* solution**
 - **Inefficiency is intolerable**

Zero Sum Game

- Total pay off to all players is the same for every instance of games
- Chess/Tic-tac-toe:
 - Win $\rightarrow 1$
 - Lose $\rightarrow -1$
 - Draw $\rightarrow 0$

Assumptions

- **Two agents acting alternately**
- **Utility values for each agent are the opposite of the other**
- Deterministic
- Fully observable
- Can generalize to stochastic games, multiple players, non zero-sum, etc
- In game theory terms:
 - “Deterministic, turn-taking, zero-sum games of perfect information”


Games as search

- Initial state: e.g. board configuration of chess
- Player: which player to give the current move
- Successor function: list of (move,state) pairs specifying legal moves.
- Terminal test: Is the game finished?
- Utility function: Gives numerical value of terminal states. E.g. win (+1), lose (-1) and draw (0) in tic-tac-toe or chess

Game Setup

- Both player wants to maximize the utility value earned at the end of the game
- We “see” the utility score as seen by MAX
- So, we can designate two players as MAX and MIN
- MAX tries to maximize its utility
- MIN tries to minimize MAX’s utility
- MAX uses search tree to determine next move.
- We play as MAX

Partial Game Tree for Tic-Tac-Toe



Size of search trees

- b = branching factor
- d = number of moves by both players
- Search tree is $O(b^d)$
- Chess
 - $b \sim 35$
 - $D \sim 100$
 - search tree is $\sim 10^{154}$ (!!)
 - completely impractical to search this
- Game-playing emphasizes being able to make optimal decisions in a finite amount of time

Optimal Strategy

An optimal strategy leads to outcomes **at least as good** as any other strategy when the opponent plays optimally


Strategy 1

The minimax algorithm


The minimax algorithm

- Find the optimal *strategy* for MAX assuming an optimal MIN opponent
- Assumption: Both players play optimally!


Two-Ply Game Tree



Two-Ply Game Tree




Two-Ply Game Tree



Two-Ply Game Tree

Minimax maximizes the utility for the worst-case outcome for max



The minimax algorithm

- Minimax value is the utility of MAX for being in the corresponding state

MINIMAX-VALUE(n) =

UTILITY(n)

If n is a terminal

$\max_{s \in \text{successors}(n)}$

MINIMAX-VALUE(s)

If n is a max node

$\min_{s \in \text{successors}(n)}$

MINIMAX-VALUE(s)

If n is a min node

Minimax algorithm

function MINIMAX-DECISION(*state*) *returns an action*

$v \leftarrow \text{MAX-VALUE}(\textit{state})$

return the *action* in SUCCESSORS(*state*) with value *v*

function MAX-VALUE(*state*) *returns a utility value*

if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow -\infty$

for *a, s* in SUCCESSORS(*state*) **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s))$

return *v*

function MIN-VALUE(*state*) *returns a utility value*

if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)


$v \leftarrow \infty$

for *a, s* in SUCCESSORS(*state*) **do**

$v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(s))$

return *v*

Minimax is done depth-first



Properties of Minimax

- Complete? Yes (if tree is finite)
- Optimal? Yes (against an optimal opponent)
- Time complexity? $O(b^m)$
- Space complexity? $O(bm)$ (depth-first exploration)
- For chess, $b \approx 35$, $m \approx 100$
→ exact solution completely infeasible

Need to speed it up.

Strategy 2

Alpha Beta Pruning

Strategy 2

Alpha Beta Pruning

We need not visit every node

Alpha-Beta Procedure

- The alpha-beta procedure can speed up a depth-first minimax search.
- **Alpha:** a lower bound on the value that a max node may ultimately be assigned

$$v \geq \alpha$$


- **Beta:** an upper bound on the value that a minimizing node may ultimately be assigned

$$v \leq \beta$$


Alpha-Beta Procedure

- At MAX node, compare the utility estimate of current node (v), with the β value, and reason whether MIN will let the game to follow this path, if not prune (i.e. $v \geq \beta$)
- At MIN node, compare the utility estimate of current node (v), with the α value, and reason whether MAX will let the game to follow this path, if not prune (i.e. $v \leq \alpha$)


α - β pruning example




α - β pruning example




α - β pruning example




α - β pruning example



α - β pruning example



α - β pruning example



Properties of α - β

- Pruning **does not** affect final result. This means that it **gets the exact same result as does full minimax**.
- Good move ordering improves effectiveness of pruning
- With "perfect ordering," time complexity = $O(b^{m/2})$
→ **doubles** depth of search

The α - β algorithm

function ALPHA-BETA-SEARCH(*state*) *returns an action*

inputs: *state*, current state in game

$v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$

return the *action* in SUCCESSORS(*state*) with value v

function MAX-VALUE(*state*, α , β) *returns a utility value*

inputs: *state*, current state in game

α , the value of the best alternative for MAX along the path to *state*

β , the value of the best alternative for MIN along the path to *state*

if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow -\infty$

for a, s in SUCCESSORS(*state*) **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s, \alpha, \beta))$

if $v \geq \beta$ **then return** v

$\alpha \leftarrow \text{MAX}(\alpha, v)$


return v

The α - β algorithm


```
function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value
  inputs: state, current state in game
            $\alpha$ , the value of the best alternative for MAX along the path to state
            $\beta$ , the value of the best alternative for MIN along the path to state

  if TERMINAL-TEST(state) then return UTILITY(state)
   $v \leftarrow +\infty$ 
  for  $a, s$  in SUCCESSORS(state) do
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(s, \alpha, \beta))$ 
    if  $v \leq \alpha$  then return  $v$ 
     $\beta \leftarrow \text{MIN}(\beta, v)$ 
  return  $v$ 
```


Example



Example



Example



What if the depth is still very large???



What if the depth is
large

Limit the search only to
a certain depth and use
a **heuristic** to provide an
estimate of the
MINIMAX value of
that node



Some heuristics for Tic-Tac-Toe

- +100 for EACH 3-in-a-line for computer.
- +10 for EACH two-in-a-line (with a empty cell) for computer.
- +1 for EACH one-in-a-line (with two empty cells) for computer.
- Negative scores for opponent, i.e., -100, -10, -1 for EACH opponent's 3-in-a-line, 2-in-a-line and 1-in-a-line.
- 0 otherwise (empty lines or lines with both computer's and opponent's seeds).

Resource

- Chapter 5
 - 5.1, 5.2, 5.3, 5.4
- Animation
 - <http://alphabetalekskamko.com/>
 - <http://homepage.ufp.pt/jtorres/ensino/ia/alfabeta.html>